

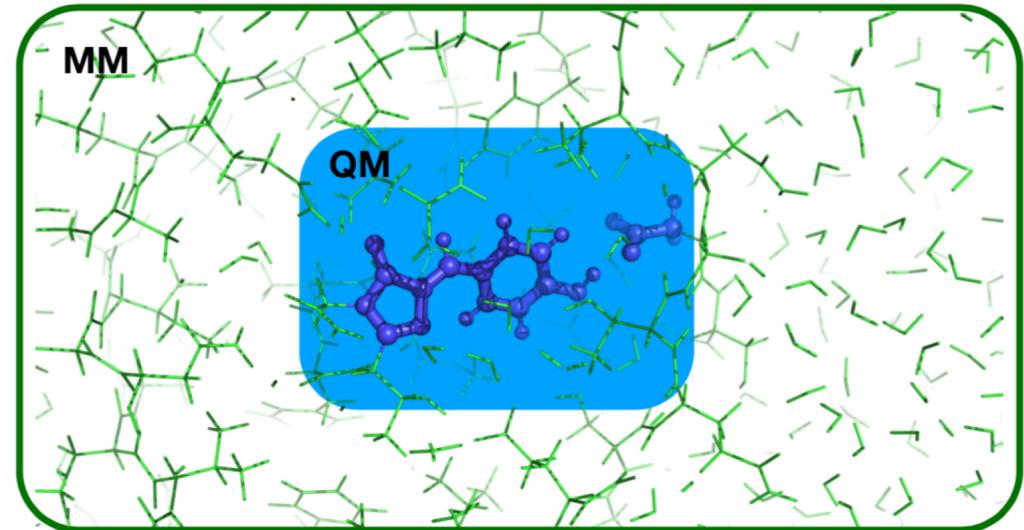
CP2K – QM/MM Practical

Aims

- Run a QM/MM simulation using CP2K on multiple nodes with MPI+OpenMP
 - Look at the effect of threading across multiple nodes
- Learn about profiling CP2K – what can we profile and why is it useful.
- Look at the communication (Message Passing) overheads and how they affect the scaling.
- Think about the different sorts of overheads that can limit scaling.

QM/MM

- MD simulation where you want the accuracy of QM in some region(s), but need classical forces (MM) in general as the system is large.
- An example of this is a protein in a liquid.
- QM energy is calculated using DFT. CP2K uses the QUICKSTEP method: a mixture of Gaussians and Plane wave basis sets.
- The coupling between the QM/MM regions is calculated using GEEP (Gaussian expansion of the Electrostatic Potential).



$$E = E_{QM} + E_{MM} + E_{QM/MM}$$

QuickStep GEEP

CP2K

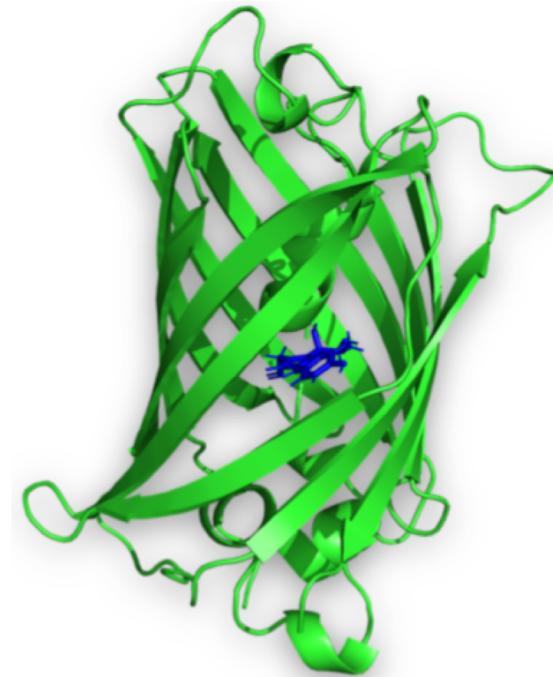
Used to perform atomistic simulations – mainly using density functional theory (DFT)

Can be run with pure MPI – **cp2k.popt**, or MPI+OpenMP - **cp2k.psmp**

Features

- Energy and Forces
- Optimisation
 - Geometry optimisation
 - Nudged elastic band
- Molecular Dynamics
 - Born-Oppenheimer MD
- Properties
 - Atomic charges (RESP, Mulliken..)
 - Spectra
 - Frequency calculations

**Green
Fluorescent
Protein with
20 QM atoms**



CP2K QM/MM best practice guide -
https://docs.bioexcel.eu/qmmm_bpg/en/main/

CP2K - Profiling

SUBROUTINE	CALLS	ASD	SELF TIME		TOTAL TIME	
			MAXIMUM	AVERAGE	MAXIMUM	AVERAGE
CP2K	1	1.0	1.012	1.093	579.165	579.166
qs_mol_dyn_low	1	2.0	0.004	0.005	566.513	566.935
qs_forces	2	3.5	0.000	0.001	465.518	465.519
qs_energies	2	4.5	0.000	0.000	457.119	457.120
scf_env_do_scf	2	5.5	0.000	0.000	453.759	453.759
scf_env_do_scf_inner_loop	90	6.2	0.002	0.006	437.466	437.466
rebuild_ks_matrix	92	8.2	0.000	0.000	365.177	365.220
qs_ks_build_kohn_sham_matrix	92	9.2	0.016	0.017	365.177	365.220
qs_ks_update_qs_env	94	7.2	0.001	0.001	356.828	356.870

- CP2K output file gives timings of the called routines – see <https://www.cp2k.org/dev:profiling>
- SELF TIME – time spent only in this routine
- TOTAL TIME – time spent in this routine, including subroutines called by it
- AVERAGE – averaged over ranks. MAXIMUM – max time of all ranks
- Difference between AVERAGE time and MAXIMUM time indicates load imbalance or synchronisation.

Exercise – Scaling of QM/MM simulations with CP2K

- Download the input files
- Create a suitable job script for running CP2K.
- Explore running the QM/MM simulation (qmmm-1.inp) on a single node of ARCHER2 with MPI+OpenMP.
- Identify the time spent in message passing routines.
- What fraction of the total time is spent in these routines?
- How does this change as the number of nodes is increased?
- Advanced: Repeat for the qmmm-4.inp system which has 77 QM atoms. How is this different?

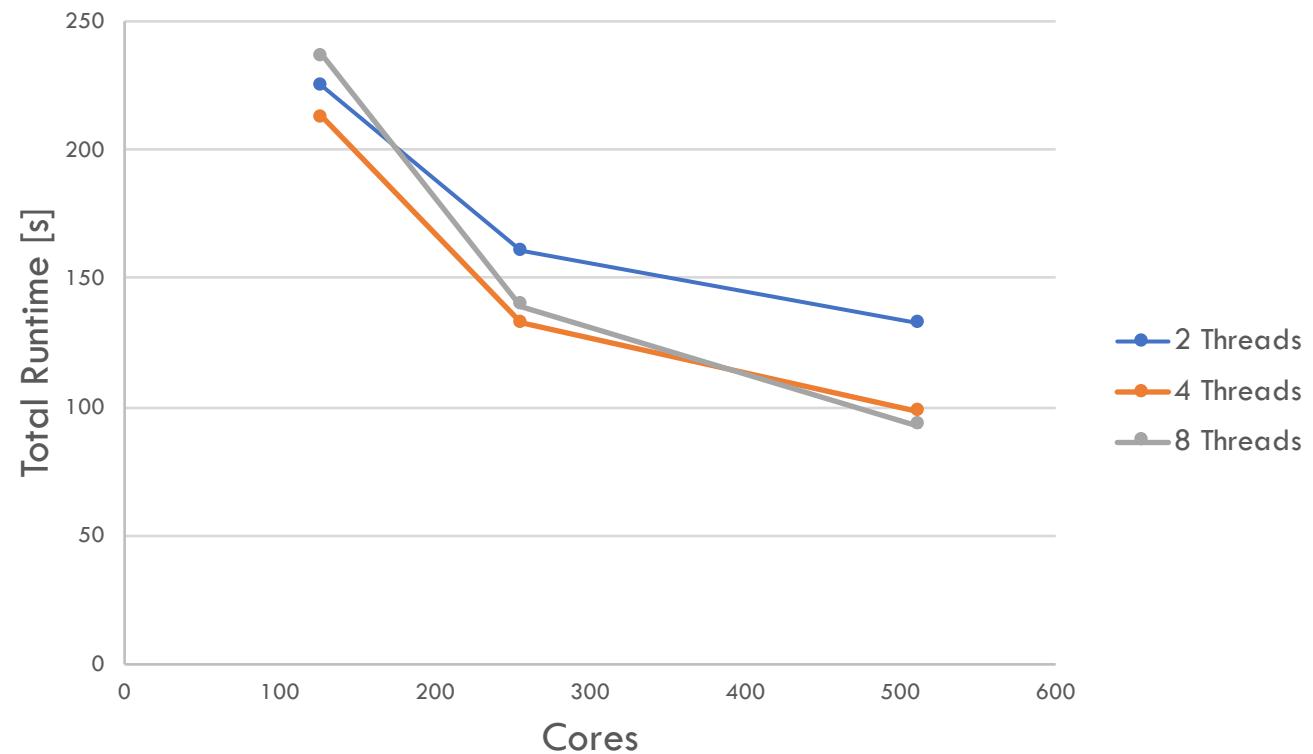
Results – Single node

- Using 1 thread results in out of memory error
 - There is not enough memory per process (128 here)
 - 2 threads allow double the memory per process, remember threads share memory so they do not need their own copy of data
- 4 threads has better performance.
- 32 threads very slow. The threads span 2 memory regions on the node which means data is slow to access

Threads	Time CP2K (s) T_total	Time mp_alltoall_z2 2v (s)	Time mp_sum_dm 3 (s)	Time mp_waitan y (s)	Total time mp_ routines (s) T_mp
1	n/a	n/a	n/a	n/a	n/a
2	224.49	18.377	14.758	4.767	37.902
4	212.471	15.599	8.821		24.42
8	236.41	20.451	5.373	7.207	33.031
16	279.018	24.936		24.819	49.755
32	379.674	35.208		26.258	61.466

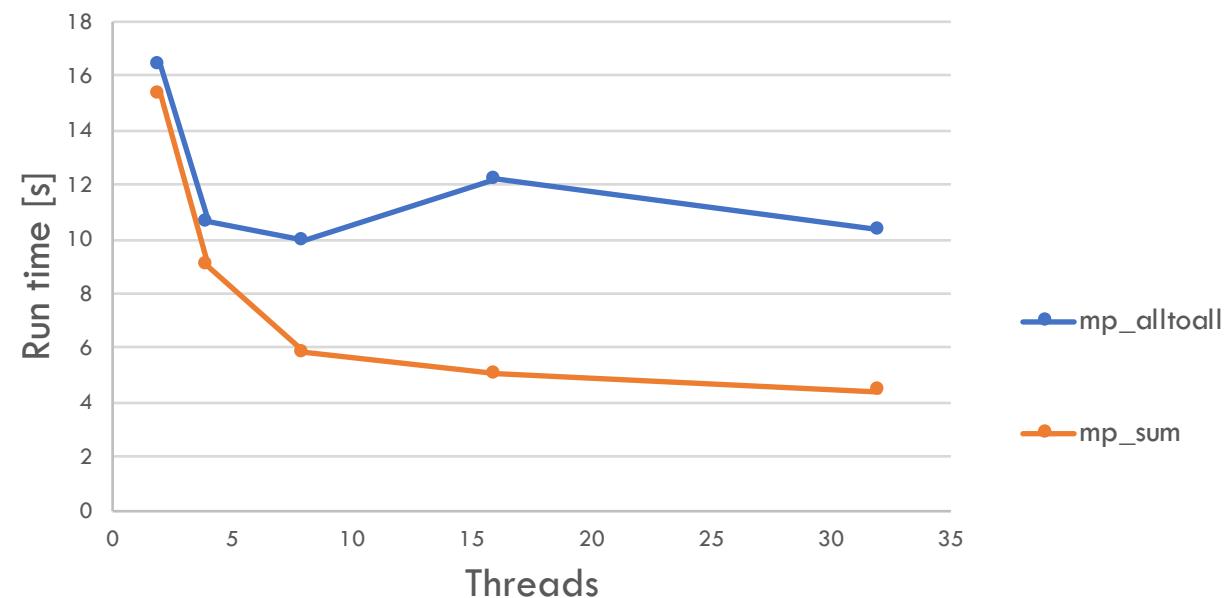
Results – Threading on multiple nodes

- Using 4 threads has the best performance on 1 and 2 nodes
 - 8 threads better on 4 nodes
 - More threads means less processes
 - Potential for less communication overhead
 - Also may exploit threading in areas of code that are not suitable for MPI
- Always a good idea to run tests to see what value is best



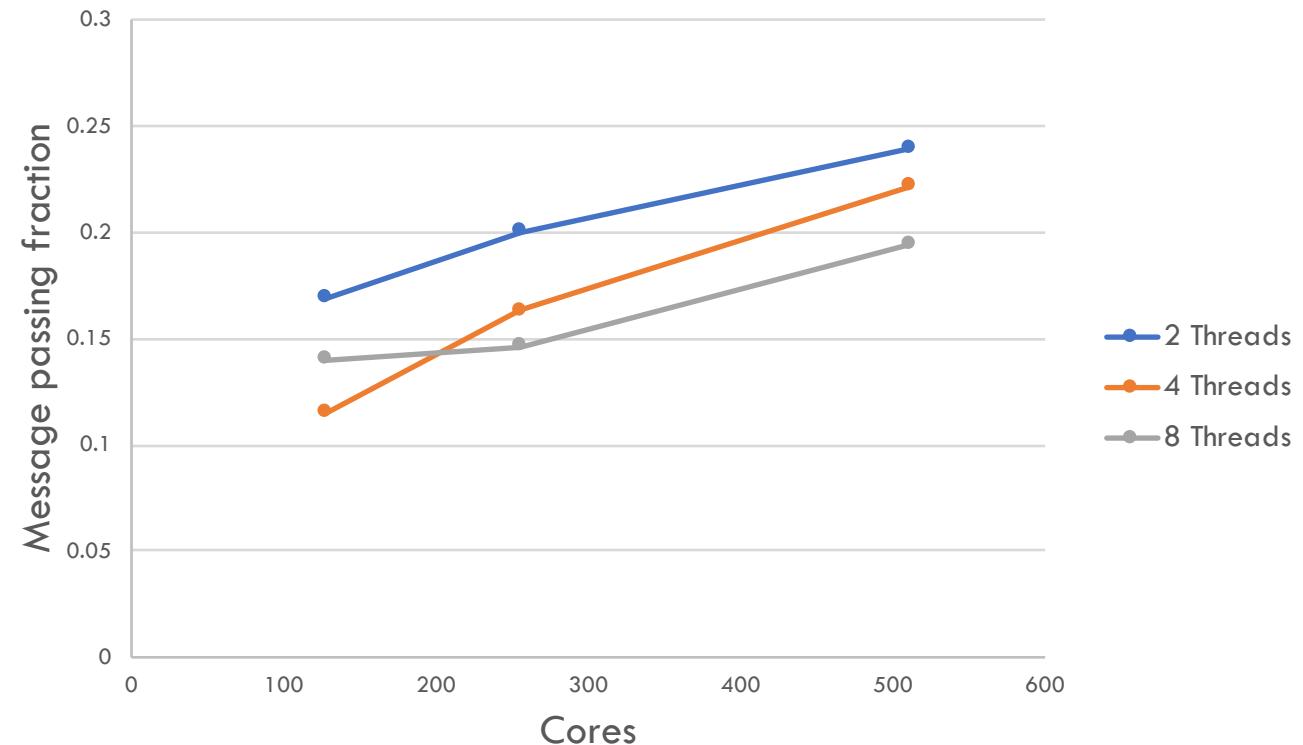
Results – Communications

- Run time of mp_sum decreases with the number of threads
 - Possibly due to less messages being sent
- Run time of mp_alltoall less clear
 - Some improvement with using threads
 - 2-4 threads look to optimal
- Hard to predict how routines and the code as a whole will be affected by changing the number of threads
 - Really have to test this



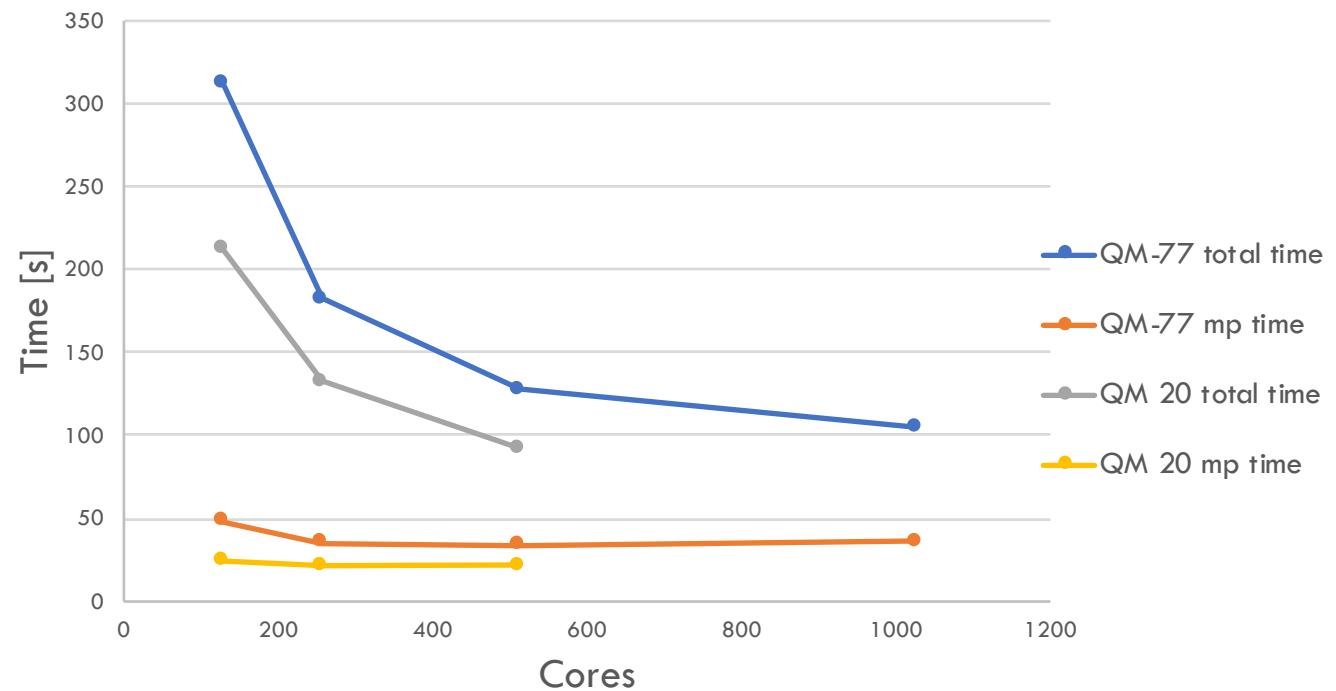
Results – Communications

- Fraction of time spent in MP routines increases with nodes
 - Other areas of the program are improving performance at a greater rate, i.e.. they have greater potential for parallel speed up.
- Changing the number of threads has the potential to reduce this fraction



Results – QM system size

- The large QM system scales better, though perhaps not as well as you would think.
- The system with the large QM region spends a greater fraction of its time doing communications.





BioExcel Partners



Ian Harrow Consulting

NORMAN
CONSULTING

Utrecht University



acrosslimits

EMBL-EBI



MAX-PLANCK-GESELLSCHAFT



MANCHESTER
1824



UNIVERSITY OF JYVÄSKYLA



JÜLICH
FORSCHUNGSZENTRUM
NBD
NOSTRUM BIODESCOVERY



INSTITUTO DE INVESTIGACIÓN BIOMÉDICA



Horizon 2020
European Union Funding
for Research & Innovation

BioExcel is funded by the European Union
Horizon 2020 program under grant
agreements 675728 and 823830.